

C1.1 Evaluation on Object Detection

Introduction

- Task: Evaluate the average precision of the data collected from SafeBench.
- Submission: the plotted P-R curve, as well as the average precision at different IoU threshold level.

Team:

Alonso Buitano Tang

Haineng Huang

Yanchen Liu

```
In [1]: import joblib
import numpy as np
import pandas as pd
from matplotlib import pyplot as plt
```

Load the Detection Results from SafeBench_v2

Link to the `results.pkl` data: <https://drive.google.com/file/d/1WyZfjhTAMbwSgkRNi1dC7r4TSf-DBil2/view?usp=sharing>

Download and extract it with this notebook.

```
In [2]: from google.colab import drive
drive.mount('/content/drive')

# OR uncomment following code, upload the dataset zip file, unzip by running the following code
! cp '/content/drive/MyDrive/24784/c1/results.zip' .
! unzip results.zip

inputs = joblib.load('results.pkl') # move the files to your current folder
inputs.keys() # dict_keys(['image_id', 'predicted_class', 'ground_truth_bbox', 'predicted_bbox', 'conf_scores', 'num_labels'])
```

Mounted at /content/drive
Archive: results.zip
inflating: results.pkl
inflating: __MACOSX/._results.pkl

```
Out[2]: dict_keys(['image_id', 'predicted_class', 'ground_truth_bbox', 'predicted_bbox', 'conf_scores', 'num_labels'])
```

```
In [3]: inputs['num_labels']
```

```
Out[3]: 92
```

Compute the IoU

```
In [4]: def box_area(box):
    """ Compute the box area, given all the vertices
    # Arguments
        box: (4, N) ndarray
    # Returns
        areas: (N, ) ndarray
    """
    # areas = (box[:,2] - box[:,0]) * (box[:,3] - box[:,1]) # TODO, compute the rectangle areas based on the vertices input
    a1, a2 = np.split(box, 2, axis=1)
    areas = (a1 - a2).prod(axis=1)
    return areas

def box_iou(box1, box2):
    """
    Compute the iou between box1 and box2
    ONLY consider the SINGLE ground truth, which is a simplified case
    box1: (N, 4) ndarray
    box2: (N, 4) ndarray
    return: (N, ) iou scores
    """
    eps = 1e-7

    a1, a2 = np.split(box1, 2, axis=1)
    b1, b2 = np.split(box2, 2, axis=1)

    c1 = np.maximum(a1, b1)
    c2 = np.minimum(a2, b2)

    inter = np.maximum(0, box_area(np.concatenate((c1,c2), axis=1)))
    union = box_area(box1) + box_area(box2) - inter

    return inter / (union + eps)
```

```
In [5]: # TODO: get the box iou scores from the function you implemented above
iou_scores = box_iou(inputs['ground_truth_bbox'], inputs['predicted_bbox'])

# Build your dataframe from the dictionary
input_dict = {
    'image_id': inputs['image_id'],
    'predicted_class': inputs['predicted_class'],
    'conf_scores': inputs['conf_scores'],
    'iou_scores': iou_scores,
}
df = pd.DataFrame.from_dict(input_dict)
df
```

```
Out[5]:
```

	image_id	predicted_class	conf_scores	iou_scores
	0	0	None	-1.000000
	1	1	car	0.671808
	2	1	car	0.563249
	3	1	car	0.452749
	4	1	car	0.428004

	1631	91	train	0.219339
	1632	91	bus	0.173378
	1633	91	person	0.101061
	1634	91	sink	0.088729
	1635	91	dining table	0.073306

1636 rows x 4 columns

```
In [6]: df = df[df.conf_scores >= 0] # drop all the frames without detection (conf scores = -1)

# get all the detection results of stop signs
df_stopsign = df.loc[(df.predicted_class=='stopsign')]
df_stopsign
```

```
Out[6]:
```

	image_id	predicted_class	conf_scores	iou_scores
	160	7	stopsign	0.071092
	345	15	stopsign	0.052572
	712	33	stopsign	0.110136
	731	34	stopsign	0.159796
	772	36	stopsign	0.128940

	1606	88	stopsign	0.119673
	1610	89	stopsign	0.997623
	1622	90	stopsign	0.995437
	1628	90	stopsign	0.052746
	1629	91	stopsign	0.995099

66 rows x 4 columns

Compute the Average Precision

```
In [7]: def interp_ap(recall, precision, method = 'interp'):
        """ Compute the average precision, given the recall and precision curves
        # Arguments
            recall: The recall curve (list)
            precision: The precision curve (list),
            methods: 'continuous', 'interp'
        # Returns
            Average precision, precision curve, recall curve
        """

        # TODO: Append sentinel values to beginning and end
        # Recall should start with 1.0 and end with 0.0
        # Precision should start with 0.0 and end with 1.0

        appended_recall = np.concatenate(([1.0], np.flip(recall), [0.0]))
        appended_prec_input = np.concatenate(([0.0], np.flip(precision), [1.0]))

        # Compute the precision envelope
        appended_prec = np.maximum.accumulate(appended_prec_input)

        appended_recall = np.flip(appended_recall)
        appended_prec = np.flip(appended_prec)
        appended_prec_input = np.flip(appended_prec_input)
        # Integrate area under curve
        if method == 'interp':
            x = np.linspace(0, 1, 101) # 101-point interp (COCO)
            ap = np.trapz(np.interp(x, appended_recall, appended_prec), x)
        else: # 'continuous', you can refer to the computation of AP in this setting when finishing the interp AP calculation
            i = np.where(appended_recall[1:] != appended_recall[:-1])[0] # points where x axis (recall) changes
            ap = np.sum((appended_recall[i + 1] - appended_recall[i]) * appended_prec[i + 1]) # area under curve

        return ap, appended_prec_input, appended_prec, appended_recall
```

```
In [8]: def compute_ap(df, num_gt, iou_thres):
        """ Compute the average precision, given the recall and precision curves
        # Arguments
            df: DataFrame inputs containing predicted classes, iou_scores, and confidence
            num_gt: The precision curve (list),
            iou_thres: IoU threshold of True Positive (TP) detection
        # Returns
            Average precision, precision curve, recall curve
        """
        df = df.sort_values('conf_scores', ascending=False)
        tp_fp = np.arange(1, df.shape[0]+1) # TODO: get the sum of all the true positive and false positive pre
        tp = ((df.iou_scores >= iou_thres) & (df.predicted_class == "stopsign")).cumsum().to_numpy() # TODO: get the true positive sets

        precision = tp / tp_fp # array (N, )
        recall = tp / num_gt # array (N, )

        ap, prec_raw, prec, recall = interp_ap(recall, precision) # get the AP and returned curve for plot

        return ap, prec_raw, prec, recall
```

Visualization






```
In [9]: def plot_pr_curves(recall, prec_raw, prec):
        """ Visualize the P-R curves
        Visualize the areas under P-R curves
        """
        plt.figure()
        plt.plot(recall, prec_raw)
        plt.plot(recall, prec)
        plt.fill_between(recall, np.zeros_like(recall), prec, color='orange', alpha=0.2)
        plt.legend(['Raw curve', 'Precision Envelope', 'Average Precision'])
        plt.title('P-R Curve')
        plt.xlabel('Recall')
        plt.ylabel('Precision')
```

Execute the AP calculation

```
In [10]: # Traverse over IoU threshold, get AP@[0.5:0.1:0.9]

for iou_thres in np.arange(0.5, 1.0, 0.1):
    ap, prec_raw, prec, recall = compute_ap(df, inputs['num_labels'], iou_thres)
    print('IoU Threshold: {:.2f}'.format(iou_thres), ' | AP@{:.2f}'.format(iou_thres), ap)
    plot_pr_curves(recall, prec_raw, prec)
```

```
IoU Threshold: 0.50 | AP@0.50 0.4272559074361477
IoU Threshold: 0.60 | AP@0.60 0.4272559074361477
IoU Threshold: 0.70 | AP@0.70 0.40686292416368364
IoU Threshold: 0.80 | AP@0.80 0.4018855902817313
IoU Threshold: 0.90 | AP@0.90 0.3466670309329718
```



IoU Calculation

Since we have the bottom-left corner of each rectangle as $a1$ and $b1 = (x1, y1)$, and the top-right corner as $a2$ and $b2 = (x2, y2)$, we obtain the maximum values of $a1$ and $b1$, giving us the larger $x1$ and $y1$ as the bottom-left corner of the intersection region. Similarly, we obtain the minimum values between $a2$ and $b2$, giving us the smaller $x2$ and $y2$ as the top-right corner of the rectangle. Using this method, if the two boxes don't intersect, the area of intersection will always be negative, so we use $\max(0, \text{box_area})$ in order to set negative values to 0 (intersection area = 0 when no intersection). Union area is calculated by summing the area of both boxes and subtracting the intersection area. From there, the IoU is obtained by dividing the intersection area over the union.

AP trend vs IoU threshold

When we increase the IoU threshold, we can see that the calculated AP value gets smaller, this makes sense since a higher threshold means that we are making a stricter threshold for true positives, meaning that we are likely to have less tp and more fp. This leads to a lower precision and recall, making the area under the P-R curve smaller. This can be observed in the P-R plots above, where at higher thresholds the precision falls faster, and the maximum value of the recall (excluding the sentinel value of 1) gets smaller.

In [10]:

